

# 10-703 Deep RL and Controls

## OpenAI Gym Recitation

Devin Schwab

Spring 2017

# Table of Contents

Introduction

Basic API

Basic Datatypes

Creating an Environment


Monitoring and Scoring

Conclusion


# What is OpenAI Gym?

- ▶ A standard Python API for RL environments
- ▶ A set of tools to **measure agent performance**
- ▶ An online scoreboard for comparing and benchmarking approaches
- ▶ <https://gym.openai.com/>


# Domain Examples




(a) Toy Text



(b) Atari



(c) Controls



(d) MuJoCo



(e) Doom



(f) Minecraft

# VirtualEnv Installation

- ▶ It is recommended that you install the gym and any dependencies in a **virtualenv**
- ▶ The following steps will create a virtualenv with the gym installed

```
virtualenv openai-gym-demo  
source openai-gym-demo/bin/activate  
pip install -U gym[all]  
python -c 'import gym; gym.make("FrozenLake-v0")'
```

# Table of Contents

Introduction

Basic API


Basic Datatypes

Creating an Environment

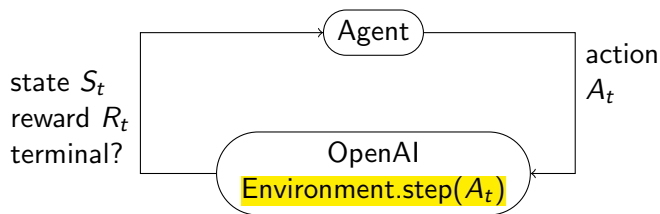
Monitoring and Scoring

Conclusion

# Basic RL Setup



# Basic RL Setup





# Basic Agent Loop

```
import gym
env = gym.make("Taxi-v2")
observation = env.reset()
for _ in range(1000):
    env.render()
    # your agent here (this takes random actions)
    action = env.action_space.sample()
    observation, reward, done, info = env.step(action)
    if done:
        env.render()
        break
```

# Creating an Instance

- ▶ Each gym environment has a unique name of the form  $([A-Za-z0-9]+-)^v([0-9]+)$
- ▶ To create an environment from the name use the  
`env = gym.make(env_name)`
- ▶ For example, to create a Taxi environment:  
`env = gym.make('Taxi-v2')`

# Reset Function

- ▶ Used to reinitialize a new episode
- ▶ Returns the initial state

```
init_state = env.reset()
```

# Step Function

```
step(action) -> (next_state,  
                 reward,  
                 is_terminal,  
                 debug_info)
```

- ▶ Performs the specified action and returns the resulting state
- ▶ The main method your agent interacts with

# Render

- ▶ Optional method
- ▶ Used to display the state of your environment
- ▶ Useful for debugging and qualitatively comparing different agent policies

# Basic Agent Demo

`demos/basic_agent.py`

# Table of Contents

Introduction

Basic API

Basic Datatypes

Creating an Environment

Monitoring and Scoring

Conclusion

# Datatypes

- ▶ **Reward** : float
- ▶ **Terminal** : bool
- ▶ **Action** : Depends on environment
- ▶ **State** : Depends on environment



# Example State Representations


	[[[0 0 0]
	[0 0 0]
	[0 0 0]
	...
	[0 0 0]
	[0 0 0]
	[0 0 0]]]
128	
(a) Taxi-v2	(b) Breakout-v0

Figure: State Representations

# Example State Representations



(a) Taxi-v2



(b) Breakout-v0

Figure: State Representations

# Example Action Representations

1	[0, 40.5, 0., -180., .5, 99.2]
(a) Taxi-v2	(b) Soccer-v0

Figure: State Representations

# Example Action Representations

1	[0, 40.5, 0., -180., .5, 99.2]
(a) Taxi-v2	(b) Soccer-v0

Figure: State Representations

How do you tell what the state and action space is for an environment?

# Environment Space Attributes

- ▶ Most environments **have two special attributes:**
  - ▶ `action_space`
  - ▶ `observation_space`
- ▶ These contain instances of `gym.spaces` classes
- ▶ **Makes it easy to find out what are valid states and actions**
- ▶ There is a convenient `sample` method to generate uniform random samples in the space.

# gym.spaces

- ▶ Action spaces and State spaces are defined by instances of classes of the gym.spaces modules
- ▶ Included types are:
  - ▶ gym.spaces.Discrete
  - ▶ gym.spaces.MultiDiscrete
  - ▶ gym.spaces.Box
  - ▶ gym.spaces.Tuple
- ▶ All instances have a sample method which will sample random instances within the space

## gym.spaces.Discrete

- ▶ The homework environments will use this type of space
- ▶ Specifies a space containing  $n$  discrete points
- ▶ Each point is mapped to an integer from  $[0, n - 1]$
- ▶ `Discrete(10)`
  - ▶ A space containing 10 items mapped to integers in  $[0, 9]$
  - ▶ `sample` will return integers such as 0, 3, and 9.

## `gym.spaces.MultiDiscrete`

- ▶ You will use this to implement an environment in the homework
- ▶ Species a space containing  $k$  dimensions each with a separate number of discrete points.
- ▶ Each point in the space is represented by a vector of integers of length  $k$
- ▶ `MultiDiscrete([(1, 3), (0, 5)])`
  - ▶ A space with  $k = 2$  dimensions
  - ▶ First dimension has 4 points mapped to integers in  $[1, 3]$
  - ▶ Second dimension has 6 points mapped to integers in  $[0, 5]$
  - ▶ `sample` will return a vector such as  $[2, 5]$  and  $[1, 3]$



## gym.spaces.Box

- ▶ Used for multidimensional continuous spaces with bounds
- ▶ You will see environments with these types of state and action spaces in future homeworks
- ▶ `Box(np.array((-1.0, -2.0)), np.array((1.0, 2.0)))`
  - ▶ A 2D continuous state space
  - ▶ First dimension has values in range  $[-1.0, 1.0]$
  - ▶ Second dimension has values in range  $[-2.0, 2.0]$
  - ▶ `sample` will return a vector such as  $[-.55, 2.]$  and  $[.768, -1.55]$

# Table of Contents

Introduction

Basic API

Basic Datatypes

Creating an Environment

Monitoring and Scoring

Conclusion

# gym.Env Class

- ▶ All environments should inherit from `gym.Env`
- ▶ At a minimum you must override a handful of methods:
  - ▶ `_step`
  - ▶ `_reset`
- ▶ At a minimum you must provide the following attributes
  - ▶ `action_space`
  - ▶ `observation_space`

# Subclass Methods

- ▶ `_step` is the same api as the `step` function used in the example
- ▶ `_reset` is the same api as the `reset` function in the example
- ▶ You may also provide the following methods for **additional functionality**:
  - ▶ `_render`
  - ▶ `_close`
  - ▶ `_configure`
  - ▶ `_seed`

# Attributes

- ▶ `observation_space` represents the state space
- ▶ `action_space` represents the action space
- ▶ Both are instances of `gym.spaces` classes
- ▶ You can also provide a `reward_range`, but this defaults to  $(-\infty, \infty)$

# Registration

- ▶ How do you get your environment to work with `gym.make()`?

# Registration

- ▶ How do you get your environment to work with `gym.make()`?
  - ▶ You must register it!

# Registration Example

```
from gym.envs.registration import register
register(
    id='Deterministic-4x4-FrozenLake-v0',
    entry_point='gym.envs.toy_text.frozen_lake:FrozenLakeEnv',
    kwargs={'map_name': '4x4',
            'is_slippery': False})
```



# Registration Example

- ▶ `id` : the environment name used with `gym.make`
- ▶ `entry_point` : module path and class name of environment
- ▶ `kwargs`: dictionary of keyword arguments to environment constructor

# Discrete Environment Class

- ▶ A subclass of the `gym.Env` which provides the following attributes
  - ▶ `nS` : number of states
  - ▶ `nA` : number of actions
  - ▶ `P` : model of environment
  - ▶ `isd` : initial state distribution

# Model

- ▶ **P is a dictionary of dictionary of lists**  
`P[s][a] == [(prob, next_state, reward, terminal), ...]`
- ▶ **isd is a list or array of length nS**  
`isd == [0., 0., 1., 0.]`

# FrozenLake-v0 Example

demos/frozen\_lake\_demo.py

# Table of Contents

Introduction

Basic API

Basic Datatypes

Creating an Environment

Monitoring and Scoring

Conclusion

# OpenAI Gym Scoreboard

- ▶ The gym also includes an **online scoreboard**
- ▶ Gym provides an API to automatically record:
  - ▶ learning curves of cumulative reward vs episode number
  - ▶ Videos of the agent executing its policy
- ▶ You can see other people's solutions and compete for the best scoreboard

# Monitor Wrapper

```
import gym
from gym import wrappers
env = gym.make('CartPole-v0')
env = wrappers.Monitor(env, '/tmp/cartpole-experiment-1')
for i_episode in range(20):
    observation = env.reset()
    for t in range(100):
        env.render()
        print(observation)
        action = env.action_space.sample()
        observation, reward, done, info = env.step(action)
        if done:
            print("Episode finished after {} timesteps".format(t+1))
            break
env.close()
gym.upload('/tmp/cartpole-experiment-1', api_key='blah')
```

# Scoreboard Demo

demos/monitor\_demo.py



# Table of Contents

Introduction

Basic API

Basic Datatypes

Creating an Environment

Monitoring and Scoring

Conclusion

# Summary

- ▶ OpenAI Gym provides a standardized API for RL environments
- ▶ Gym also provides an online scoreboard for sharing and comparing results/techniques
- ▶ With only a few functions you can have your own gym environment to use with your RL algorithms

Thank You

Questions